

Going deeper with dplyr: New features in 0.3 and 0.4

Introduction

In August 2014, I created a 40-minute video tutorial introducing the key functionality of the dplyr package in R, using dplyr version 0.2. Since then, there have been two significant updates to dplyr (0.3 and 0.4), introducing a ton of new features.

This document (created in March 2015) covers the most useful new features in 0.3 and 0.4, as well as other functionality that I didn't cover last time (though it is not necessarily new). My new video tutorial walks through the code below in detail.

If you have not watched the previous tutorial, I recommend you do so first since it covers some dplyr basics that will not be covered in this tutorial.

Loading dplyr and the nycflights13 dataset

Although my last tutorial used data from the hflights package, Hadley Wickham has rewritten the dplyr vignettes to use the nycflights13 package instead, and so I'm also using nycflights13 for the sake of consistency.

```
# remove flights data if you just finished my previous tutorial
rm(flights)
```

```
# load packages
suppressMessages(library(dplyr))
#Make sure to install.packages("nycflights13")
library(nycflights13)
```

```
# print the flights dataset from nycflights13
flights
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Choosing columns: select, rename

```
# besides just using select() to pick columns...
```

```
flights %>% select(carrier, flight)
```

```
## # A tibble: 336,776 × 2
```

```
##   carrier flight
```

```
##   <chr>   <int>
```

```
## 1     UA    1545
```

```
## 2     UA    1714
```

```
## 3     AA    1141
```

```
## 4     B6     725
```

```
## 5     DL     461
```

```
## 6     UA    1696
```

```
## 7     B6     507
```

```
## 8     EV    5708
```

```
## 9     B6      79
```

```
## 10    AA     301
```

```
## # ... with 336,766 more rows
```

```
# ...you can use the minus sign to hide columns
```

```
flights %>% select(-month, -day)
```

```
## # A tibble: 336,776 × 17
```

```
##   year dep_time sched_dep_time dep_delay arr_time sched_arr_time
```

```
##   <int>   <int>         <int>      <dbl>   <int>         <int>
```

```
## 1   2013     517           515         2       830           819
```

```
## 2   2013     533           529         4       850           830
```

```
## 3   2013     542           540         2       923           850
```

```
## 4   2013     544           545        -1      1004          1022
```

```
## 5   2013     554           600        -6       812           837
```

```
## 6   2013     554           558        -4       740           728
```

```
## 7   2013     555           600        -5       913           854
```

```
## 8   2013     557           600        -3       709           723
```

```
## 9   2013     557           600        -3       838           846
```

```
## 10  2013     558           600        -2       753           745
```

```
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
```

```
## #   time_hour <dtm>
```

```
# hide a range of columns
```

```
flights %>% select(-(dep_time:arr_delay))
```

```
# hide any column with a matching name
```

```
flights %>% select(-contains("time"))
```

```
# pick columns using a character vector of column names
```

```
cols <- c("carrier", "flight", "tailnum")
```

```
flights %>% select(one_of(cols))
```

```
## # A tibble: 336,776 × 3
```

```
##   carrier flight tailnum
```

```
##   <chr>   <int>   <chr>
```

```
## 1     UA    1545 N14228
```

```
## 2     UA    1714 N24211
```

```
## 3      AA    1141  N619AA
## 4      B6     725  N804JB
## 5      DL     461  N668DN
## 6      UA    1696  N39463
## 7      B6     507  N516JB
## 8      EV    5708  N829AS
## 9      B6      79  N593JB
## 10     AA     301  N3ALAA
## # ... with 336,766 more rows
```

```
# select() can be used to rename columns, though all columns not mentioned are dropped
flights %>% select(tail = tailnum)
```

```
## # A tibble: 336,776 × 1
##   tail
##   <chr>
## 1 N14228
## 2 N24211
## 3 N619AA
## 4 N804JB
## 5 N668DN
## 6 N39463
## 7 N516JB
## 8 N829AS
## 9 N593JB
## 10 N3ALAA
## # ... with 336,766 more rows
```

```
# rename() does the same thing, except all columns not mentioned are kept
flights %>% rename(tail = tailnum)
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tail <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Choosing rows: filter, between, slice, sample_n, top_n, distinct

```
# filter() supports the use of multiple conditions
flights %>% filter(dep_time >= 600, dep_time <= 605)
```

```
## # A tibble: 2,460 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     600           600         0     851
## 2  2013     1     1     600           600         0     837
## 3  2013     1     1     601           600         1     844
## 4  2013     1     1     602           610        -8     812
## 5  2013     1     1     602           605        -3     821
## 6  2013     1     2     600           600         0     814
## 7  2013     1     2     600           605        -5     751
## 8  2013     1     2     600           600         0     819
## 9  2013     1     2     600           600         0     846
## 10 2013     1     2     600           600         0     737
## # ... with 2,450 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>

# between() is a concise alternative for determining if numeric values fall in a range
flights %>% filter(between(dep_time, 600, 605))

# side note: is.na() can also be useful when filtering
flights %>% filter(!is.na(dep_time))

# slice() filters rows by position
flights %>% slice(1000:1005)
```

```
## # A tibble: 6 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     2     809           810        -1     950
## 2  2013     1     2     810           800        10    1008
## 3  2013     1     2     811           815        -4    1100
## 4  2013     1     2     811           815        -4    1126
## 5  2013     1     2     811           820        -9     944
## 6  2013     1     2     815           815         0    1109
## # ... with 12 more variables: sched_arr_time <int>, arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>

# keep the first three rows within each group
flights %>% group_by(month, day) %>% slice(1:3)
```

```
## Source: local data frame [1,095 × 19]
## Groups: month, day [365]
##
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     2      42          2359        43     518
## 5  2013     1     2    126          2250       156     233
## 6  2013     1     2    458           500        -2     703
## 7  2013     1     3     32          2359        33     504
```

```
## 8 2013 1 3 50 2145 185 203
## 9 2013 1 3 235 2359 156 700
## 10 2013 1 4 25 2359 26 505
## # ... with 1,085 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# sample three rows from each group
flights %>% group_by(month, day) %>% sample_n(3)
```

```
## Source: local data frame [1,095 x 19]
```

```
## Groups: month, day [365]
```

```
##
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
```

```
## 1 2013     1     1    2053           2055         -2    2254
```

```
## 2 2013     1     1    1558           1359        119    1718
```

```
## 3 2013     1     1    2209           2155         14    2400
```

```
## 4 2013     1     2    1902           1905         -3    2203
```

```
## 5 2013     1     2    1257           1300         -3    1529
```

```
## 6 2013     1     2    2051           1929         82    2256
```

```
## 7 2013     1     3    2046           2045          1    2340
```

```
## 8 2013     1     3    1455           1446          9    1801
```

```
## 9 2013     1     3    1510           1510          0    1736
```

```
## 10 2013     1     4    1643           1615         28    1857
```

```
## # ... with 1,085 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
# keep three rows from each group with the top dep_delay
```

```
flights %>% group_by(month, day) %>% top_n(3, dep_delay)
```

```
## Source: local data frame [1,108 x 19]
```

```
## Groups: month, day [365]
```

```
##
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
```

```
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
```

```
## 1 2013     1     1     848           1835        853    1001
```

```
## 2 2013     1     1    1815           1325        290    2120
```

```
## 3 2013     1     1    2343           1724        379     314
```

```
## 4 2013     1     2    1412           838         334    1710
```

```
## 5 2013     1     2    1607          1030        337    2003
```

```
## 6 2013     1     2    2131          1512        379    2340
```

```
## 7 2013     1     3    2008          1540        268    2339
```

```
## 8 2013     1     3    2012          1600        252    2314
```

```
## 9 2013     1     3    2056          1605        291    2239
```

```
## 10 2013     1     4    2058          1730        208      2
```

```
## # ... with 1,098 more rows, and 12 more variables: sched_arr_time <int>,
```

```
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
```

```
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
```

```
## #   minute <dbl>, time_hour <dtm>
```

```
# also sort by dep_delay within each group
```

```
flights %>% group_by(month, day) %>% top_n(3, dep_delay) %>% arrange(desc(dep_delay))
```

```
## Source: local data frame [1,108 x 19]
## Groups: month, day [365]
##
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     9     641           900        1301    1242
## 2  2013     6    15    1432          1935        1137    1607
## 3  2013     1    10    1121          1635        1126    1239
## 4  2013     9    20    1139          1845        1014    1457
## 5  2013     7    22     845          1600        1005    1044
## 6  2013     4    10    1100          1900         960    1342
## 7  2013     3    17    2321           810         911     135
## 8  2013     6    27     959          1900         899    1236
## 9  2013     7    22    2257           759         898     121
## 10 2013    12     5     756          1700         896    1058
## # ... with 1,098 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

```
# unique rows can be identified using unique() from base R
flights %>% select(origin, dest) %>% unique()
```

```
## # A tibble: 224 × 2
##   origin dest
##   <chr> <chr>
## 1   EWR   IAH
## 2   LGA   IAH
## 3   JFK   MIA
## 4   JFK   BQN
## 5   LGA   ATL
## 6   EWR   ORD
## 7   EWR   FLL
## 8   LGA   IAD
## 9   JFK   MCO
## 10  LGA   ORD
## # ... with 214 more rows
```

```
# dplyr provides an alternative that is more "efficient"
flights %>% select(origin, dest) %>% distinct()
```

```
# side note: when chaining, you don't have to include the parentheses if there are no arguments
flights %>% select(origin, dest) %>% distinct
```

Adding new variables: mutate, transmute, add_rownames

```
# mutate() creates a new variable (and keeps all existing variables)
flights %>% mutate(speed = distance/air_time*60)
```

```
## # A tibble: 336,776 × 20
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
```

```
## 3 2013 1 1 542 540 2 923
## 4 2013 1 1 544 545 -1 1004
## 5 2013 1 1 554 600 -6 812
## 6 2013 1 1 554 558 -4 740
## 7 2013 1 1 555 600 -5 913
## 8 2013 1 1 557 600 -3 709
## 9 2013 1 1 557 600 -3 838
## 10 2013 1 1 558 600 -2 753
## # ... with 336,766 more rows, and 13 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>, speed <dbl>
```

```
# transmute() only keeps the new variables
flights %>% transmute(speed = distance/air_time*60)
```

```
## # A tibble: 336,776 × 1
##   speed
##   <dbl>
## 1 370.0441
## 2 374.2731
## 3 408.3750
## 4 516.7213
## 5 394.1379
## 6 287.6000
## 7 404.4304
## 8 259.2453
## 9 404.5714
## 10 318.6957
## # ... with 336,766 more rows
```

```
# example data frame with row names
mtcars %>% head()
```

```
##           mpg cyl disp  hp drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110 3.90 2.620 16.46  0  1   4    4
## Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02  0  1   4    4
## Datsun 710     22.8   4  108  93 3.85 2.320 18.61  1  1   4    1
## Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44  1  0   3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0   3    2
## Valiant        18.1   6  225 105 2.76 3.460 20.22  1  0   3    1
```

```
# add_rownames() turns row names into an explicit variable
mtcars %>% add_rownames("model") %>% head()
```

```
## Warning: Deprecated, use tibble::rownames_to_column() instead.
```

```
## # A tibble: 6 × 12
##   model      mpg   cyl  disp    hp  drat    wt   qsec    vs  am
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Mazda RX4      21.0     6   160   110   3.90  2.620 16.46     0    1
## 2 Mazda RX4 Wag  21.0     6   160   110   3.90  2.875 17.02     0    1
## 3 Datsun 710     22.8     4   108    93   3.85  2.320 18.61     1    1
## 4 Hornet 4 Drive  21.4     6   258   110   3.08  3.215 19.44     1    0
## 5 Hornet Sportabout 18.7     8   360   175   3.15  3.440 17.02     0    0
## 6 Valiant        18.1     6   225   105   2.76  3.460 20.22     1    0
## # ... with 2 more variables: gear <dbl>, carb <dbl>
```

```
# side note: dplyr no longer prints row names (ever) for local data frames
mtcars %>% tbl_df()
```

```
## # A tibble: 32 × 11
##   mpg   cyl  disp    hp  drat    wt  qsec    vs  am  gear  carb
## *   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  21.0     6 160.0   110  3.90  2.620 16.46    0    1     4     4
## 2  21.0     6 160.0   110  3.90  2.875 17.02    0    1     4     4
## 3  22.8     4 108.0    93  3.85  2.320 18.61    1    1     4     1
## 4  21.4     6 258.0   110  3.08  3.215 19.44    1    0     3     1
## 5  18.7     8 360.0   175  3.15  3.440 17.02    0    0     3     2
## 6  18.1     6 225.0   105  2.76  3.460 20.22    1    0     3     1
## 7  14.3     8 360.0   245  3.21  3.570 15.84    0    0     3     4
## 8  24.4     4 146.7    62  3.69  3.190 20.00    1    0     4     2
## 9  22.8     4 140.8    95  3.92  3.150 22.90    1    0     4     2
## 10 19.2     6 167.6   123  3.92  3.440 18.30    1    0     4     4
## # ... with 22 more rows
```

Grouping and counting: summarise, tally, count, group_size, n_groups, ungroup

```
# summarise() can be used to count the number of rows in each group
flights %>% group_by(month) %>% summarise(cnt = n())
```

```
## # A tibble: 12 × 2
##   month  cnt
##   <int> <int>
## 1     1 27004
## 2     2 24951
## 3     3 28834
## 4     4 28330
## 5     5 28796
## 6     6 28243
## 7     7 29425
## 8     8 29327
## 9     9 27574
## 10    10 28889
## 11    11 27268
## 12    12 28135
```

```
# tally() and count() can do this more concisely
flights %>% group_by(month) %>% tally()
flights %>% count(month)
```

```
# you can sort by the count
flights %>% group_by(month) %>% summarise(cnt = n()) %>% arrange(desc(cnt))
```

```
## # A tibble: 12 × 2
##   month  cnt
##   <int> <int>
## 1     7 29425
## 2     8 29327
## 3    10 28889
```



```
## 4      3 28834
## 5      5 28796
## 6      4 28330
## 7      6 28243
## 8     12 28135
## 9      9 27574
## 10     11 27268
## 11      1 27004
## 12      2 24951
```

```
# tally() and count() have a sort parameter for this purpose
```

```
flights %>% group_by(month) %>% tally(sort=TRUE)
flights %>% count(month, sort=TRUE)
```

```
# you can sum over a specific variable instead of simply counting rows
```

```
flights %>% group_by(month) %>% summarise(dist = sum(distance))
```

```
## # A tibble: 12 × 2
```

```
##   month   dist
##   <int>   <dbl>
## 1     1 27188805
## 2     2 24975509
## 3     3 29179636
## 4     4 29427294
## 5     5 29974128
## 6     6 29856388
## 7     7 31149199
## 8     8 31149334
## 9     9 28711426
## 10    10 30012086
## 11    11 28639718
## 12    12 29954084
```

```
# tally() and count() have a wt parameter for this purpose
```

```
flights %>% group_by(month) %>% tally(wt = distance)
flights %>% count(month, wt = distance)
```

```
# group_size() returns the counts as a vector
```

```
flights %>% group_by(month) %>% group_size()
```

```
## [1] 27004 24951 28834 28330 28796 28243 29425 29327 27574 28889 27268
## [12] 28135
```

```
# n_groups() simply reports the number of groups
```

```
flights %>% group_by(month) %>% n_groups()
```

```
## [1] 12
```

```
# group by two variables, summarise, arrange (output is possibly confusing)
```

```
flights %>% group_by(month, day) %>% summarise(cnt = n()) %>% arrange(desc(cnt)) %>% print(n = 40)
```

```
## Source: local data frame [365 x 3]
```

```
## Groups: month [12]
```

```
##
```

```
##   month   day   cnt
##   <int> <int> <int>
## 1     11    27  1014
```

```
## 2      7      11 1006
## 3      7       8 1004
## 4      7      10 1004
## 5     12       2 1004
## 6      7      18 1003
## 7      7      25 1003
## 8      7      12 1002
## 9      7       9 1001
## 10     7      17 1001
## 11     7      31 1001
## 12     8       7 1001
## 13     8       8 1001
## 14     8      12 1001
## 15     7      22 1000
## 16     7      24 1000
## 17     8       1 1000
## 18     8       5 1000
## 19     8      15 1000
## 20    11      21 1000
## 21     7      15  999
## 22     7      19  999
## 23     7      26  999
## 24     7      29  999
## 25     8       2  999
## 26     8       9  999
## 27    11      22  999
## 28     8      16  998
## 29     7      23  997
## 30     7      30  997
## 31     8      14  997
## 32     7      16  996
## 33     8       6  996
## 34     8      19  996
## 35     9      13  996
## 36     9      26  996
## 37     9      27  996
## 38     4      15  995
## 39     6      20  995
## 40     6      26  995
## # ... with 325 more rows
```

```
# ungroup() before arranging to arrange across all groups
flights %>% group_by(month, day) %>% summarise(cnt = n()) %>% ungroup() %>% arrange(desc(cnt))
```

```
## # A tibble: 365 × 3
##   month   day   cnt
##   <int> <int> <int>
## 1     11    27  1014
## 2      7     11  1006
## 3      7      8  1004
## 4      7     10  1004
## 5     12      2  1004
## 6      7     18  1003
## 7      7     25  1003
## 8      7     12  1002
```

```
## 9      7      9 1001
## 10     7     17 1001
## # ... with 355 more rows
```

Creating data frames: data_frame

`data_frame()` is a better way than `data.frame()` for creating data frames. Benefits of `data_frame()`:

- You can use previously defined columns to compute new columns.
- It never coerces column types.
- It never munges column names.
- It never adds row names.
- It only recycles length 1 input.
- It returns a local data frame (a `tbl_df`).

data_frame() example

```
data_frame(a = 1:6, b = a*2, c = 'string', 'd+e' = 1) %>% glimpse()
```

```
## Observations: 6
## Variables: 4
## $ a    <int> 1, 2, 3, 4, 5, 6
## $ b    <dbl> 2, 4, 6, 8, 10, 12
## $ c    <chr> "string", "string", "string", "string", "string", "string"
## $ d+e  <dbl> 1, 1, 1, 1, 1, 1
```

data.frame() example

```
data.frame(a = 1:6, c = 'string', 'd+e' = 1) %>% glimpse()
```

```
## Observations: 6
## Variables: 3
## $ a    <int> 1, 2, 3, 4, 5, 6
## $ c    <fctr> string, string, string, string, string, string
## $ d.e  <dbl> 1, 1, 1, 1, 1, 1
```

Joining (merging) tables: left_join, right_join, inner_join, full_join, semi_join, anti_join

create two simple data frames

```
(a <- data_frame(color = c("green","yellow","red"), num = 1:3))
```

```
## # A tibble: 3 × 2
##   color  num
##   <chr> <int>
## 1 green     1
## 2 yellow    2
## 3  red      3
```

```
(b <- data_frame(color = c("green","yellow","pink"), size = c("S","M","L")))
```

```
## # A tibble: 3 × 2
##   color size
##   <chr> <chr>
## 1 green  S
## 2 yellow M
## 3 pink   L
```

```
# only include observations found in both "a" and "b" (automatically joins on variables that appear in
inner_join(a, b)
```

```
## Joining, by = "color"
```

```
## # A tibble: 2 × 3
##   color  num size
##   <chr> <int> <chr>
## 1 green     1    S
## 2 yellow    2    M
```

```
# include observations found in either "a" or "b"
full_join(a, b)
```

```
## Joining, by = "color"
```

```
## # A tibble: 4 × 3
##   color  num size
##   <chr> <int> <chr>
## 1 green     1    S
## 2 yellow    2    M
## 3 red       3 <NA>
## 4 pink     NA    L
```

```
# include all observations found in "a"
left_join(a, b)
```

```
## Joining, by = "color"
```

```
## # A tibble: 3 × 3
##   color  num size
##   <chr> <int> <chr>
## 1 green     1    S
## 2 yellow    2    M
## 3 red       3 <NA>
```

```
# include all observations found in "b"
right_join(a, b)
```

```
## Joining, by = "color"
```

```
## # A tibble: 3 × 3
##   color  num size
##   <chr> <int> <chr>
## 1 green     1    S
## 2 yellow    2    M
## 3 pink     NA    L
```

```
# right_join(a, b) is identical to left_join(b, a) except for column ordering
left_join(b, a)
```

```
## Joining, by = "color"
```

```
## # A tibble: 3 × 3
##   color size  num
##   <chr> <chr> <int>
## 1 green    S     1
## 2 yellow   M     2
## 3 pink     L    NA
```

```
# filter "a" to only show observations that match "b"
semi_join(a, b)

## Joining, by = "color"

## # A tibble: 2 × 2
##   color  num
##   <chr> <int>
## 1 green    1
## 2 yellow   2

# filter "a" to only show observations that don't match "b"
anti_join(a, b)
```

```
## Joining, by = "color"

## # A tibble: 1 × 2
##   color  num
##   <chr> <int>
## 1 red    3

# sometimes matching variables don't have identical names
b <- b %>% rename(col = color)

# specify that the join should occur by matching "color" in "a" with "col" in "b"
inner_join(a, b, by=c("color" = "col"))

## # A tibble: 2 × 3
##   color  num size
##   <chr> <int> <chr>
## 1 green    1 S
## 2 yellow   2 M
```

Viewing more output: print, View

```
# specify that you want to see more rows
flights %>% print(n = 15)

## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517             515           2     830
## 2  2013     1     1     533             529           4     850
## 3  2013     1     1     542             540           2     923
## 4  2013     1     1     544             545          -1    1004
## 5  2013     1     1     554             600          -6     812
## 6  2013     1     1     554             558          -4     740
## 7  2013     1     1     555             600          -5     913
## 8  2013     1     1     557             600          -3     709
## 9  2013     1     1     557             600          -3     838
## 10 2013     1     1     558             600          -2     753
## 11 2013     1     1     558             600          -2     849
## 12 2013     1     1     558             600          -2     853
## 13 2013     1     1     558             600          -2     924
## 14 2013     1     1     558             600          -2     923
```

```
## 15 2013      1      1      559          600      -1      941
## # ... with 3.368e+05 more rows, and 12 more variables:
## #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# specify that you want to see ALL rows (don't run this!)
flights %>% print(n = Inf)
```

```
# specify that you want to see all columns
flights %>% print(width = Inf)
```

```
## # A tibble: 336,776 × 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>
## 1  2013     1     1     517           515         2     830
## 2  2013     1     1     533           529         4     850
## 3  2013     1     1     542           540         2     923
## 4  2013     1     1     544           545        -1    1004
## 5  2013     1     1     554           600        -6     812
## 6  2013     1     1     554           558        -4     740
## 7  2013     1     1     555           600        -5     913
## 8  2013     1     1     557           600        -3     709
## 9  2013     1     1     557           600        -3     838
## 10 2013     1     1     558           600        -2     753
##   sched_arr_time arr_delay carrier flight tailnum origin dest air_time
##             <int>      <dbl>   <chr>  <int>   <chr>   <chr> <chr>   <dbl>
## 1             819         11     UA    1545  N14228   EWR   IAH     227
## 2             830         20     UA    1714  N24211   LGA   IAH     227
## 3             850         33     AA    1141  N619AA   JFK   MIA     160
## 4            1022        -18     B6     725  N804JB   JFK   BQN     183
## 5             837        -25     DL     461  N668DN   LGA   ATL     116
## 6             728         12     UA    1696  N39463   EWR   ORD     150
## 7             854         19     B6     507  N516JB   EWR   FLL     158
## 8             723        -14     EV    5708  N829AS   LGA   IAD      53
## 9             846         -8     B6      79  N593JB   JFK   MCO     140
## 10            745          8     AA     301  N3ALAA   LGA   ORD     138
##   distance  hour minute      time_hour
##       <dbl> <dbl>  <dbl>         <dtm>
## 1     1400     5     15 2013-01-01 05:00:00
## 2     1416     5     29 2013-01-01 05:00:00
## 3     1089     5     40 2013-01-01 05:00:00
## 4     1576     5     45 2013-01-01 05:00:00
## 5       762     6      0 2013-01-01 06:00:00
## 6       719     5     58 2013-01-01 05:00:00
## 7     1065     6      0 2013-01-01 06:00:00
## 8       229     6      0 2013-01-01 06:00:00
## 9       944     6      0 2013-01-01 06:00:00
## 10      733     6      0 2013-01-01 06:00:00
## # ... with 336,766 more rows
```

```
# show up to 1000 rows and all columns
flights %>% View()
```

```
# set option to see all columns and fewer rows
```

```
options(dplyr.width = Inf, dplyr.print_min = 6)

# reset options (or just close R)
options(dplyr.width = NULL, dplyr.print_min = 10)
```

Resources

- Release announcements for version 0.3 and version 0.4
- dplyr reference manual and vignettes
- Two-table vignette covering joins and set operations
- RStudio's Data Wrangling Cheat Sheet for dplyr and tidyr
- dplyr GitHub repo and list of releases

Data School

- Blog
- Email newsletter
- YouTube channel

< END OF DOCUMENT >